

LangChain Demo Document

A sample PDF for testing LangChain document loaders

Introduction

This document was generated to demonstrate loading PDF files with LangChain. It contains structured text content that can be parsed, chunked, and used in retrieval-augmented generation (RAG) pipelines.

Key Concepts

- Document Loaders read files (PDF, TXT, CSV, HTML) into LangChain Document objects.
- Each Document has **page_content** (the text) and **metadata** (source, page number, etc.).
- Text Splitters break large documents into smaller chunks for embedding and retrieval.
- Vector stores index these chunks so a retriever can find the most relevant ones for a query.

Sample Code

```
from langchain_community.document_loaders import PyPDFLoader

loader = PyPDFLoader("docs/langchain_demo.pdf")
documents = loader.load()
print(f"Loaded {len(documents)} pages")
for doc in documents:
    print(doc.page_content[:200])
    print(doc.metadata)
```

Why This Matters

PDFs are one of the most common document formats in the world. Being able to reliably extract text from them is a foundational skill for building AI applications that can reason over real-world documents. LangChain makes this straightforward with its loader ecosystem.

Conclusion

You now have a properly formatted PDF to test with. Try loading it with PyPDFLoader, experiment with different chunking strategies, and build a simple RAG pipeline on top of it. Happy building!